

# Assignment 1

Name: Pasupulety Chethan Krishna Venkat

Roll Number: 21CS30036

```
# import all the necessary libraries here
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import numpy as np
from graphviz import Digraph
import matplotlib.pyplot as plt

df = pd.read_csv('../dataset/decision-tree.csv')
print(df.shape)

# split the dataset into train and test and validation here
train_df, temp_data_df = train_test_split(df, test_size=0.5,
random_state=42)
validation_df, test_df = train_test_split(temp_data_df, test_size=0.4,
random_state=42)

# Convert DataFrame to numpy arrays
train_data = train_df.to_numpy()
test_data = test_df.to_numpy()
validation_data = validation_df.to_numpy()

# Split data into X and y
X_train = train_data[:, :-1]           # all rows, all the
features and no labels                 # all rows, label
y_train = train_data[:, -1]           only

X_test = test_data[:, :-1]
y_test = test_data[:, -1]

X_validation = validation_data[:, :-1]
y_validation = validation_data[:, -1]

# a list to store the names of the features
feature_names = list(df.columns[:-1])

# Using the id3 algorithm to build the decision tree
```

```
# Since the data is continous,we need to find the best attribute to  
split the data and best threshold value for that attribute  
# We will use the information gain to find the best attribute and  
threshold value  
# We will use the entropy function to calculate the information gain
```

```
# Entropy function
```

```
def entropy(y):  
    entropy = 0  
    unique, counts = np.unique(y, return_counts=True)  
    for i in range(len(unique)):  
        entropy += -(counts[i]/len(y))*np.log2(counts[i]/len(y))  
    return entropy
```

```
# Information gain function
```

```
def info_gain(X, y, attribute_index, threshold):  
    entropy_before = entropy(y)  
    entropy_after = 0  
    y_left = []  
    y_right = []  
    for i in range(len(y)):  
        if X[i][attribute_index] <= threshold:  
            y_left.append(y[i])  
        else:  
            y_right.append(y[i])  
    entropy_after = (len(y_left)/len(y))*entropy(y_left) +  
    (len(y_right)/len(y))*entropy(y_right)  
    return entropy_before - entropy_after
```

```
# Finding the best attribute and threshold value
```

```
def best_attribute_threshold(X, y):  
  
    best_attribute_index = 0  
    best_threshold = 0  
    max_info_gain = 0  
    for i in range(X.shape[1]):  
        for j in range(X.shape[0]):  
            info_gain_val = info_gain(X, y, i, X[j][i])  
            if info_gain_val > max_info_gain:  
                max_info_gain = info_gain_val  
                best_attribute_index = i  
                best_threshold = X[j][i]  
    return (best_attribute_index, best_threshold)
```

```
# Node class for the decision tree
```

```
class Node:  
    def __init__(self, attribute_index=None, threshold=None,  
left=None, right=None, label=None):
```

```

        self.attribute_index = attribute_index
        self.threshold = threshold
        self.left = left
        self.right = right
        self.label = label

    def is_leaf_node(self):
        return self.label is not None

# Building the decision tree using the ID3 algorithm
# tree having min_size as stopping criteria

def build_tree(X, y, min_size):

    if len(y) <= min_size:
        unique, counts = np.unique(y, return_counts=True)
        return Node(label=unique[np.argmax(counts)])
    best_attribute_index, best_threshold = best_attribute_threshold(X,
y)
    y_left = []
    y_right = []
    X_left = []
    X_right = []
    for i in range(len(y)):
        if X[i][best_attribute_index] <= best_threshold:
            y_left.append(y[i])
            X_left.append(X[i])
        else:
            y_right.append(y[i])
            X_right.append(X[i])
    if len(y_left) == 0 or len(y_right) == 0:
        unique, counts = np.unique(y, return_counts=True)
        return Node(label=unique[np.argmax(counts)])
    left = build_tree(np.array(X_left), np.array(y_left), min_size)
    right = build_tree(np.array(X_right), np.array(y_right), min_size)
    return Node(best_attribute_index, best_threshold, left, right)

# Predicting the label for a single data point
def predict(node, data_point):

    if node.is_leaf_node():
        return node.label
    if data_point[node.attribute_index] <= node.threshold:
        return predict(node.left, data_point)

```

```

        else:
            return predict(node.right, data_point)

# Predicting the labels for a set of data points
def predict_labels(root, X):

    y_pred = []
    for i in range(len(X)):
        y_pred.append(predict(root, X[i]))
    return np.array(y_pred)

# Calculating the accuracy of the model
def accuracy(y_pred, y_true):

    correct = 0
    for i in range(len(y_pred)):
        if y_pred[i] == y_true[i]:
            correct += 1
    return correct/len(y_pred)

# printing the tree with feature names and threshold value using graphviz and saving it as png
def visualize_tree(node, feature_names, graph=None):
    if graph is None:
        graph = Digraph(format='png') # You can change the format if you prefer a different image format

    if node.is_leaf_node():
        graph.node(str(id(node)), label=str(node.label))
    else:
        feature_name = feature_names[node.attribute_index]
        graph.node(str(id(node)), label=f"{feature_name}\nThreshold {node.threshold}")
        if node.left:
            visualize_tree(node.left, feature_names, graph)
            graph.edge(str(id(node)), str(id(node.left)),
label='True')
        if node.right:
            visualize_tree(node.right, feature_names, graph)
            graph.edge(str(id(node)), str(id(node.right)),
label='False')

    return graph

# pruning with respect to the validation set and depth of the tree
def reduced_error_pruning(node, X_validation, y_validation):

```

```

# Write your code to perform the reduced error pruning algorithm
# Return the root node of the pruned decision tree
# YOUR CODE HERE
if node.is_leaf_node():
    return node
if node.left.is_leaf_node() and node.right.is_leaf_node():
    y_pred = predict_labels(node, X_validation)
    accuracy_before = accuracy(y_pred, y_validation)

    # changing the nodes label to the most common label in the
validation set
    unique, counts = np.unique(y_validation, return_counts=True)
    node.label = unique[np.argmax(counts)]

    y_pred = predict_labels(node, X_validation)
    accuracy_after = accuracy(y_pred, y_validation)
    if accuracy_after >= accuracy_before:
        node.left = None
        node.right = None
        return node
    else:
        node.label = None
        return node
    node.left = reduced_error_pruning(node.left, X_validation,
y_validation)
    node.right = reduced_error_pruning(node.right, X_validation,
y_validation)
    return node

# pruning with respect to the validation set and depth of the tree
# changing depth of the tree
def build_tree_depth(X, y, min_size, depth):
    if len(y) <= min_size:
        unique, counts = np.unique(y, return_counts=True)
        return Node(label=unique[np.argmax(counts)])
    best_attribute_index, best_threshold = best_attribute_threshold(X,
y)
    y_left = []
    y_right = []
    X_left = []
    X_right = []
    for i in range(len(y)):
        if X[i][best_attribute_index] <= best_threshold:
            y_left.append(y[i])
            X_left.append(X[i])
        else:

```

```

        y_right.append(y[i])
        X_right.append(X[i])
    if len(y_left) == 0 or len(y_right) == 0:
        unique, counts = np.unique(y, return_counts=True)
        return Node(label=unique[np.argmax(counts)])
    if depth == 0:
        unique, counts = np.unique(y, return_counts=True)
        return Node(label=unique[np.argmax(counts)])
    left = build_tree_depth(np.array(X_left), np.array(y_left),
min_size, depth-1)
    right = build_tree_depth(np.array(X_right), np.array(y_right),
min_size, depth-1)
    return Node(best_attribute_index, best_threshold, left, right)

```

*# plotting accuracy for the graph for different depth values*

```

def plot_accuracy(X_train, y_train, X_test, y_test, X_validation,
y_validation):
    train_accuracy = []
    test_accuracy = []
    validation_accuracy = []
    for i in range(1, 11):
        root = build_tree_depth(X_train, y_train, 10, i)
        y_pred = predict_labels(root, X_train)
        train_accuracy.append(accuracy(y_pred, y_train))
        y_pred = predict_labels(root, X_test)
        test_accuracy.append(accuracy(y_pred, y_test))
        y_pred = predict_labels(root, X_validation)
        validation_accuracy.append(accuracy(y_pred, y_validation))
    return train_accuracy, test_accuracy, validation_accuracy

```

*# Main function*

```

def main():
    # Build the tree
    root = build_tree(X_train, y_train, 10)

```

```

# Visualize the tree
graph = visualize_tree(root, feature_names)
graph.render('decision_tree', view=True)

# Predict the labels for the test data
y_pred = predict_labels(root, X_test)

# Calculate the accuracy of the model on train and test data
train_accuracy = accuracy(predict_labels(root, X_train), y_train)
test_accuracy = accuracy(y_pred, y_test)

print(f"Train accuracy: {train_accuracy}")
print(f"Test accuracy: {test_accuracy}")

# plot accuracy vs depth plots

train_accuracy, test_accuracy, validation_accuracy =
plot_accuracy(X_train, y_train, X_test, y_test, X_validation,
y_validation)

# Create a range of depths (assuming you have 10 depths)
depths = list(range(1, 11))

# Plot the accuracy values
plt.figure(figsize=(10, 6))
plt.plot(depths, train_accuracy, label='Train Accuracy')
plt.plot(depths, test_accuracy, label='Test Accuracy')
plt.plot(depths, validation_accuracy, label='Validation Accuracy')
plt.xlabel('Tree Depth')
plt.ylabel('Accuracy')
plt.title('Accuracy vs. Tree Depth')
plt.legend()
plt.grid(True)
plt.show()

if __name__ == '__main__':

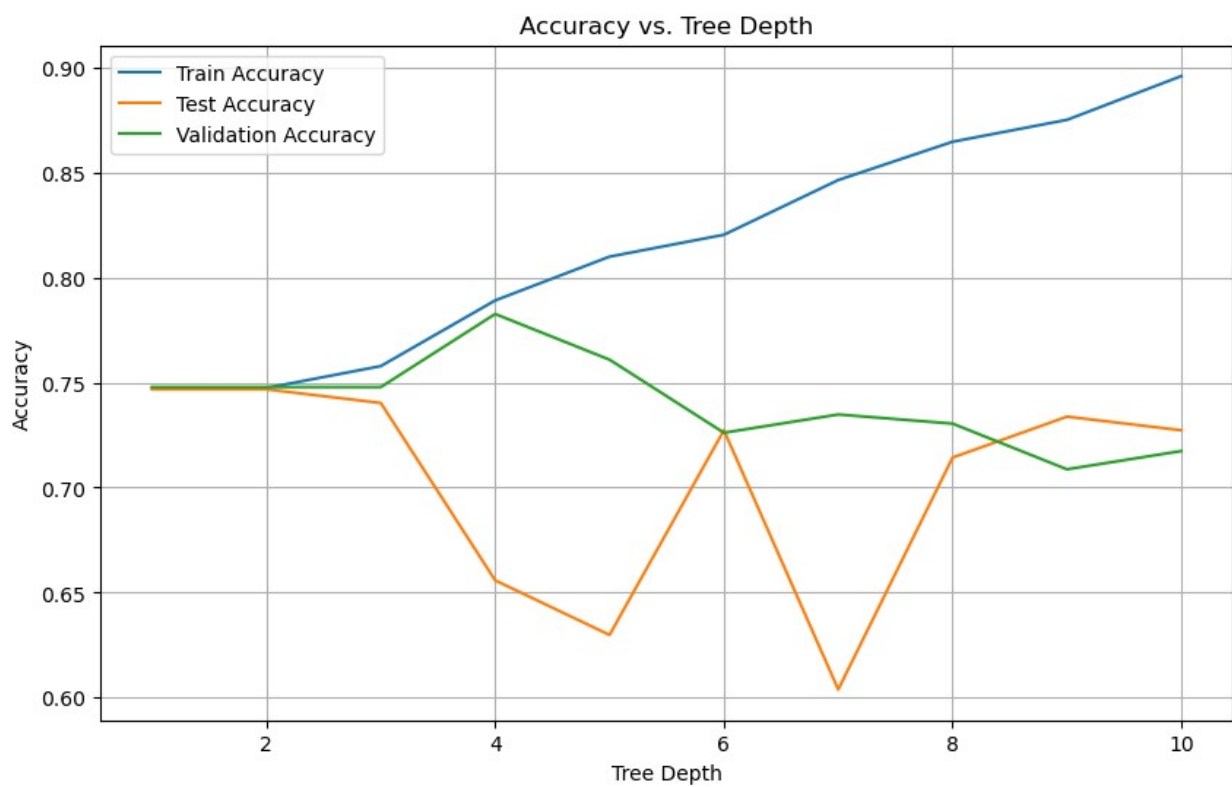
```

```
main()
```

```
(768, 9)
```

```
Train accuracy: 0.9322916666666666
```

```
Test accuracy: 0.7012987012987013
```



```
# Main function
def main():
    # Build the tree
```



```

root = build_tree(X_train, y_train, 10)

# Visualize the tree
graph = visualize_tree(root, feature_names)
graph.render('decision_tree', view=True)

# Predict the labels for the test data
y_pred = predict_labels(root, X_test)

# Calculate the accuracy of the model on train and test data
train_accuracy = accuracy(predict_labels(root, X_train), y_train)
test_accuracy = accuracy(y_pred, y_test)

print(f"Train accuracy: {train_accuracy}")
print(f"Test accuracy: {test_accuracy}")

# plot accuracy vs depth plots

train_accuracy, test_accuracy, validation_accuracy =
plot_accuracy(X_train, y_train, X_test, y_test, X_validation,
y_validation)

# Create a range of depths (assuming you have 10 depths)
depths = list(range(1, 11))

# Plot the accuracy values
plt.figure(figsize=(10, 6))
plt.plot(depths, train_accuracy, label='Train Accuracy')
plt.plot(depths, test_accuracy, label='Test Accuracy')
plt.plot(depths, validation_accuracy, label='Validation Accuracy')
plt.xlabel('Tree Depth')
plt.ylabel('Accuracy')
plt.title('Accuracy vs. Tree Depth')
plt.legend()
plt.grid(True)
plt.show()

```

```
if __name__ == '__main__':  
    main()
```